

## Exercise1

### Make a PMF

The GSS dataset has been pre-loaded for you into a DataFrame called `gss`. You can explore it in the IPython Shell to get familiar with it.

In this exercise, you'll focus on one variable in this dataset, 'year', which represents the year each respondent was interviewed.

The `Pmf` class you saw in the video has already been created for you. You can access it outside of DataCamp via the `empiricaldist` library.

### Instructions

- Make a PMF for year with `normalize=False` and display the result.

In [ ]:

```
# Compute the PMF for year
pmf_year = Pmf(____, ____=____)

# Print the result
print(pmf_year)

#_____#
#Solutions

# Compute the PMF for year
pmf_year = Pmf(gss['year'], normalize=False)

# Print the result
print(pmf_year)
```

## Exercise2

### Plot a PMF

Now let's plot a PMF for the age of the respondents in the GSS dataset. The variable 'age' contains respondents' age in years.

### Instructions

- Select the 'age' column from the `gss` DataFrame and store the result in `age`.
- Make a normalized PMF of age. Store the result in `pmf_age`.
- Plot `pmf_age` as a bar chart.

In [ ]:

```
# Select the age column
age = ____

# Make a PMF of age
pmf_age = ____

# Plot the PMF

# Label the axes

# _____#
#Solutions

# Select the age column
age = gss['age']

# Make a PMF of age
pmf_age = Pmf(age)

# Plot the PMF
pmf_age.bar()

# Label the axes
plt.xlabel('Age')
plt.ylabel('PMF')
plt.show()
```

## Exercise3

# Make a CDF

In this exercise, you'll make a CDF and use it to determine the fraction of respondents in the GSS dataset who are OLDER than 30.

The GSS dataset has been preloaded for you into a DataFrame called gss.

As with the Pmf class from the previous lesson, the Cdf class you just saw in the video has been created for you, and you can access it outside of DataCamp via the empiricaldist library.

## Instructions

- Select the 'age' column. Store the result in age
- Compute the CDF of age. Store the result in cdf\_age.
- Calculate the CDF of 30.

In [ ]:

```
# Select the age column  
# Compute the CDF of age  
# Calculate the CDF of 30  
# _____#  
#Solutions  
  
# Select the age column  
age = gss['age']  
  
# Compute the CDF of age  
cdf_age = Cdf(age)  
  
# Calculate the CDF of 30  
print(cdf_age(30))
```

## Exercise4

# Compute IQR

Recall from the video that the interquartile range (IQR) is the difference between the 75th and 25th percentiles. It is a measure of variability that is robust in the presence of errors or extreme values.

In this exercise, you'll compute the interquartile range of income in the GSS dataset. Income is stored in the 'realinc' column, and the CDF of income has already been computed and stored in cdf\_income.

## Instructions

- Calculate the 75th percentile of income and store it in percentile\_75th.
- Calculate the 25th percentile of income and store it in percentile\_25th.
- Calculate the interquartile range of income. Store the result in iqr.

In [ ]:

```
# Calculate the 75th percentile
percentile_75th = ____

# Calculate the 25th percentile

# Calculate the interquartile range

# Print the interquartile range

# _____#
#Solutions

# Calculate the 75th percentile
percentile_75th = cdf_income.inverse(0.75)

# Calculate the 25th percentile
percentile_25th = cdf_income.inverse(0.25)

# Calculate the interquartile range
iqr = percentile_75th - percentile_25th

# Print the interquartile range
print(iqr)
```

## Exercise5

### Plot a CDF

The distribution of income in almost every country is long-tailed; that is, there are a small number of people with very high incomes.

In the GSS dataset, the variable 'realinc' represents total household income, converted to 1986 dollars. We can get a sense of the shape of this distribution by plotting the CDF.

### Instructions

- Select 'realinc' from the gss dataset.
- Make a Cdf object called cdf\_income.
- Create a plot of cdf\_income using .plot().

In [ ]:

```
# Select realinc
income = ____

# Make the CDF
cdf_income = ____

# Plot it
____

# Label the axes
plt.xlabel('Income (1986 USD)')
plt.ylabel('CDF')
plt.show()

# _____#
#Solutions

# Select realinc
income = gss['realinc']

# Make the CDF
cdf_income = Cdf(income)

# Plot it
cdf_income.plot()

# Label the axes
plt.xlabel('Income (1986 USD)')
plt.ylabel('CDF')
plt.show()
```

## Exercise6

### Extract education levels

Let's create Boolean Series to identify respondents with different levels of education.

In the U.S, 12 years of education usually means the respondent has completed high school (secondary education). A respondent with 14 years of education has probably completed an associate degree (two years of college); someone with 16 years has probably completed a bachelor's degree (four years of college).

### Instructions

- Complete the line that identifies respondents with associate degrees, that is, people with 14 or more years of education but less than 16.
- Complete the line that identifies respondents with 12 or fewer years of education.
- Confirm that the mean of high is the fraction we computed in the previous exercise, about 53%.

In [ ]:

```

# Select educ
educ = gss['educ']

# Bachelor's degree
bach = (educ >= 16)

# Associate degree
assc = _____

# High school (12 or fewer years of education)
high = _____
print(high.mean())

# _____#
#Solutions

# Select educ
educ = gss['educ']

# Bachelor's degree
bach = (educ >= 16)

# Associate degree
assc = (educ >= 14) & (educ < 16)

# High school
high = (educ <= 12)
print(high.mean())

```

## Exercise7

### Plot income CDFs

Let's now see what the distribution of income looks like for people with different education levels. You can do this by plotting the CDFs. Recall how Allen plotted the income CDFs of respondents interviewed before and after 1995:

```

Cdf(income[pre95]).plot(label='Before 1995')
Cdf(income[~pre95]).plot(label='After 1995')

```

You can assume that Boolean Series have been defined, as in the previous exercise, to identify respondents with different education levels: high, assc, and bach.

### Instructions

- Fill in the missing lines of code to plot the CDFs.

In [ ]:

```

income = gss['realinc']

# Plot the CDFs
____(label='High school')
____(label='Associate')
____(label='Bachelor')

# Label the axes
plt.xlabel('Income (1986 USD)')
plt.ylabel('CDF')
plt.legend()
plt.show()

# _____#
#Solutions

income = gss['realinc']

# Plot the CDFs
Cdf(income[high]).plot(label='High school')
Cdf(income[assc]).plot(label='Associate')
Cdf(income[bach]).plot(label='Bachelor')

# Label the axes
plt.xlabel('Income (1986 USD)')
plt.ylabel('CDF')
plt.legend()
plt.show()

```

## Exercise8

# Distribution of income

In many datasets, the distribution of income is approximately lognormal, which means that the logarithms of the incomes fit a normal distribution. We'll see whether that's true for the GSS data. As a first step, you'll compute the mean and standard deviation of the log of incomes using NumPy's `np.log10()` function.

Then, you'll use the computed mean and standard deviation to make a norm object using the `scipy.stats.norm()` function.

## Instructions

- Extract 'realinc' from gss and compute its logarithm using `np.log10()`.
- Compute the mean and standard deviation of the result.
- Make a norm object by passing the computed mean and standard deviation to `norm()`.

In [ ]:

```

# Extract realinc and compute its log
income = gss['realinc']
log_income = ____

# Compute mean and standard deviation
mean = ____
std = ____
print(mean, std)

# Make a norm object
from scipy.stats import norm
dist = ____

#_____#
#Solutions

# Extract realinc and compute its log
income = gss['realinc']
log_income = np.log10(income)

# Compute mean and standard deviation
mean = log_income.mean()
std = log_income.std()
print(mean, std)

# Make a norm object
from scipy.stats import norm
dist = norm(mean, std)

```

## Exercise9

# Comparing CDFs

To see whether the distribution of income is well modeled by a lognormal distribution, we'll compare the CDF of the logarithm of the data to a normal distribution with the same mean and standard deviation. These variables from the previous exercise are available for use:

```

# Extract realinc and compute its log
log_income = np.log10(gss['realinc'])

# Compute mean and standard deviation
mean, std = log_income.mean(), log_income.std()

# Make a norm object
from scipy.stats import norm
dist = norm(mean, std)

```

`dist` is a `scipy.stats.norm` object with the same mean and standard deviation as the data. It provides `.cdf()`, which evaluates the normal cumulative distribution function.

Be careful with capitalization: `Cdf()`, with an uppercase `C`, creates `Cdf` objects. `dist.cdf()`, with a lowercase `c`, evaluates the normal cumulative distribution function.



## Instructions

- Evaluate the normal cumulative distribution function using `dist.cdf`.
- Use the `Cdf()` function to compute the CDF of `log_income`.
- Plot the result.

In [ ]:

```
# Evaluate the model CDF
xs = np.linspace(2, 5.5)
ys = _____

# Plot the model CDF
plt.clf()
plt.plot(xs, ys, color='gray')

# Create and plot the Cdf of log_income
_____

# Label the axes
plt.xlabel('log10 of realinc')
plt.ylabel('CDF')
plt.show()

# _____ #
#Solutions

# Evaluate the model CDF
xs = np.linspace(2, 5.5)
ys = dist.cdf(xs)

# Plot the model CDF
plt.clf()
plt.plot(xs, ys, color='gray')

# Create and plot the Cdf of Log_income
Cdf(log_income).plot()

# Label the axes
plt.xlabel('log10 of realinc')
plt.ylabel('CDF')
plt.show()
```

## Exercise10

## Comparing PDFs

In the previous exercise, we used CDFs to see if the distribution of income is lognormal. We can make the same comparison using a PDF and KDE. That's what you'll do in this exercise!

As before, the norm object `dist` is available in your workspace:

```
from scipy.stats import norm
dist = norm(mean, std)
```

Just as all norm objects have a `.cdf()` method, they also have a `.pdf()` method.

To create a KDE plot, you can use Seaborn's `kdeplot()` function. Here, Seaborn has been imported for you as `sns`.

## Instructions

- Evaluate the normal PDF using `dist`, which is a norm object with the same mean and standard deviation as the data.
- Make a KDE plot of the logarithms of the incomes, using `log_income`, which is a Series object.

In [ ]:

```
# Evaluate the normal PDF
xs = np.linspace(2, 5.5)
ys = _____

# Plot the model PDF
plt.clf()
plt.plot(xs, ys, color='gray')

# Plot the data KDE
_____

# Label the axes
plt.xlabel('log10 of realinc')
plt.ylabel('PDF')
plt.show()

# _____#
#Solutions

# Evaluate the normal PDF
xs = np.linspace(2, 5.5)
ys = dist.pdf(xs)

# Plot the model PDF
plt.clf()
plt.plot(xs, ys, color='gray')

# Plot the data KDE
sns.kdeplot(log_income)

# Label the axes
plt.xlabel('log10 of realinc')
plt.ylabel('PDF')
plt.show()
```